

• Losses and Classifiers

- $f(x; W) = Wx + b$ with three viewpoints
 - * Geometric: weights W are hyper-planes cutting up data space.
 - * Algebraic: optimization problem, try to minimize loss value.
 - * Visual: each weight corresponds to a "template" (i.e. pattern).
- Classifier loss functions : SVM (linear) vs Softmax (logistic)
 - * SVM (hinge) loss : $\sum_{j \neq y_i} \max(0, s_i - s_{y_i} + 1)$, square to penalize large errors.
 - * Softmax (cross entropy) loss : $-\log(\sum_i (e^{s_i} / \sum_j e^{s_j}))$, computes log-prob for each correct class.
 - * Full : loss + regularized term (sum of squares of all weights, prevent overfitting)
- Loss optimization : gradient descent $w' = w - \text{learningrate} \times \nabla_W L$, compute $\nabla_W L = \frac{\partial}{\partial w} L$
- Numeric gradient : $\nabla_W L = \frac{L(W+\delta) - L(W)}{\delta}$ for small δ , easy but inaccurate.
- Analytic gradient : derive ∇L as a function of W , uses computational graph data structure.
- Stacking layers : apply non-linearity since 2 linear layers can be combined into one. Feature transform of non-linear layers can map non-linear space into linear ones, easier to predict from.

• Normalizations, Initializations and Optimizers

- Normalizations : try to keep data on the same range.
 - * Batch norm : store mean-std of dataset on training (moving average), norm across each channel.
 - * Instance norm : Gaussian individual norm across channel and sample.
 - * Layer norm : Gaussian norm across each sample.
- Initialization : optimal weights to start with
 - * Normal (Gaussian) initialization tends to zero for larger (deeper) layers.
 - * Xavier initialization scales weight by $\text{std} = 1/\sqrt{D_{in}}$, still bad with ReLU.
 - * Kaiming initialization corrects ReLU by $\text{std} = \sqrt{2/D_{in}}$
- Optimizers : Update the weights to lower loss values.
 - * SGD : compute gradient descent on either full (expensive) or minibatch (approximate)
 - * Problems with SGD : unstable from high condition number, prone to local minima
 - * Momentum : $v' = v - \rho \nabla w$, $w' = w - \alpha v'$. Computes velocity from ∇w , ρ acts as friction.
 - * Nesterov Momentum : compute gradient after applying velocity (look ahead)
 - * AdaGrad : scale gradients with root-sum-of-squares of historical gradients, norm gradient across each dimension. Decays eventually.
 - * RMSProp : "Leaky AdaGrad" stores historical sum with moving average.
 - * ADAM : Momentum + RMSProp + Bias Correction (avoid explosion from initial guesses for the betas).
- DL uses first degree optimization, since second degree approx (though more accurate) has extremely high complexity.
- Learning Rate Decay : slow down training to find global minima (cosine vs linear).
- Linear warmup : increase LR for the first few epochs before decaying (start with low LR to prevent explosion)
- Dropout : neglect some features at training time, add back those features and scale accordingly on testing time.

• Convolutional Neural Networks

- Parameters : feature size, kernel size, zero-padding, stride
- LeNet5 : old architecture with $2 \times 5 \times 5$ convolution layers and fully connected classifier.
- AlexNet : first use of ReLU, heavy data augmentation, uses normalization (not common anymore). Most params are in FC layers.
- VGG : deeper network, uses smaller kernel 3×3 (combine to increase receptive field), still retains FC classifiers.
- ResNet : deeper models (100+ layers) should technically be as good as shallow ones but identity maps are difficult to optimize. Uses residual connection to emulate identity function. No more FC classifiers except very last layer.

• Object Detection

- Modal (only visible part) vs Amodal (including hidden parts) object detection.
- IoU : intersection over union of bounding box, good because it scales with box sizes.
- Multi-task training : optimize bounding box and classification losses simultaneously.
- Dealing with multiple objects - Two Stage Detector
 - * Sliding window : literally predicts every region - super slow and memory intensive
 - * Region proposal : use traditional CV to estimate 2000 regions that may contain an object.
 - * R-CNN : re-scale each region to 224×224 and feed each one to a CNN + classifier + bbox regressor. Slow since CNN is used 2000 times.
 - * Fast R-CNN : run the entire image through a CNN first.
 - * Faster R-CNN : use another CNN to compute region proposal.
- Dealing with multiple objects - One Stage Detector
 - * General idea : divides image into grid cell with assumption each grid can contain a box, background is now a class.
 - * RetinaNet : center boxes around grid, predict box with a class and scaling factor for predefined anchors (preset of aspect ratios), deal with class imbalance using focal loss (cross entropy, weighted on hard misclassified samples)
 - * FCOS : bbox are predicted as distances from center, centerness loss to ensure boxes stay close to grid. "Anchor Free", more generalizable.
- Dealing with scale - Feature Pyramid

- * Many stages of CNN at each resolution (higher stage = bigger box, lower res).
- * Add up features of higher res with upscaled lower res image.
- * Connected final features to an object detector at each res.
- NMS : iteratively pick a box and eliminating boxes with high IoU. Repeat until no box is left.
- mAP : sorts predicted score, for each score plot a point on the P-R curve, AP = AUC of this curve, mAP = average AP across classes.
- Semantic Segmentation : FCN (CNN+up/down-sampling+TransposedCNN) takes in local+global features, is end-to-end, and scale-able by size.
- Instance Segmentation : Mask-RCNN predicts a semantic mask out of each predicted bounding box.
- **Recurrent Neural Networks and Transformers**
 - Key concept : hidden state $x_t = f(h_{t-1}, x)$ where $h_t = \tanh(W_{hh}h_{t-1} + W_{hx}x_t + b_h)$ and output $y_t = W_{hy}h_t + b_y$.
 - Seq2Seq : Many to one (encodes input seq into one vector) + One to many (decodes output seq into many vectors). Start from $\langle \text{START} \rangle$ token and end with an $\langle \text{END} \rangle$ token.
 - CV use cases in image captioning (use CNN's encoded vector as another trainable term for hidden state).
 - Seq2Seq with RNN and attention : RNN generates one context vector that gets bottle-necked and reused.
 - * Takes in Q, K, V and outputs a mapped version of Q with context from V .
 - * K has the same length as V but is in the same dim as Q .
 - * Compute atten weight $W = \text{softmax}(QK^T/\sqrt{d_k})$
 - * Output is in the form WV and will have the same length as Q with dimension of V .
 - * Here S_{t+1} is from $Q = S_t$, $K = V = (h_1, h_2, \dots, h_T)$. Allows for related words to pop-up in translation (or image regions).
 - Vision Transformers : 3 ideas
 - * Attention in between CNN : Still a CNN, marginal improvements
 - * Replace CNN with local attention : tricky and still marginal improvements
 - * Pixel-wise Transformer : super memory consuming
 - * ViT : split image into patches, treat encoded vectors from each patch as sequence, add a learnable classification token at the end.
- **Generative Models**
 - Explicit/Probabilistic models : directly compute $p(x)$ over data x
 - * Autoregressive : predicts pixels using data from previous pixels (like an RNN), is extremely slow.
 - * VAEs : normal AE but latent vector z is sampled from μ and σ , allows for interpolation. Generate blurry images since L_2 is bad.
 - * DDPM : add noises to images until converges to $N(0, 1)$, U-Net takes in timestep and noisy image to predict added noises.
 - Implicit models : compute $p(x|y)$ over data x and condition y (either true vs fake or classes)
 - * GANs : trains a minimax optimization between generator and discriminator. Can mode collapse if G and D are outbalanced.
 - Stylization : pattern generation uses gram matrix (correlation matrix between each channels from a pretrained image classifier), captures sets of features (thus a texture).
 - Style transfer takes in content and style images. Minimizes gram matrix loss with style and feature loss with content. A dedicated CNN can be trained for a style to improve on speed.
 - Explainable AI Stuff: try to show what's going on in a CNN
 - * Nearest neighbor or dim reduction (i.e. PCA or t-SNE) on output vectors.
 - * Masks part of input image to show impact on prediction. Saliency via backprop / Guided backprop computes gradient of prediction in respect to pixel.
 - * Gradient Ascent : Start from zero image and generates an ideal image of each class. Penalizes L_2 norm of generated images.
 - * Adversarial Attacks : generate fake samples to fool the network - easy
 - * Adversarial Defance : make your model smart enough to not be fooled - hard
 - * DeepDream : amply certain features on already existing images, generates traces of that feature.
- **Self-supervised learning**
 - Sparse AE : reconstruct input from a bottleneck latent vector.
 - Denoising AE : removed added noise into clean images.
 - Context prediction : predict relative location of two patches from an image. Find neighboring patches with KNN.
 - Context encoder : paint missing pixels from a mask. Uses an AE and randomly masks parts of an image.
 - Colorization : model must learn underlying image structure to know which color to fill.
 - Deep Clustering : generated pseudo label with random model and K-means, re-train model with cluster assignments.
 - RotNet : predict image rotations using rotation as pretext
 - ExemplarCNN : predict where an augmented image is from, given a pool of input images.
 - Contrastive learning : minimize similarity across classes, minimize similarity within the same class.
 - Masked AE : reconstruct with most patches are missing with ViT before using a Seq2Seq decoder (another ViT) to generate missing parts.
 - Multi-modal SSL : takes in other inputs (i.e. text caption) i.e. CLIP - matching text and images.